

# PES University, Bangalore

Established under Karnataka Act No. 16 of 2013

UE20CS312 - Data Analytics - Worksheet 3b - AR and MA models

Ishita Bharadwaj - PES1UG20CS648

Collaborated with Hita Juneja - PES1UG20CS645

## AR and MA models

Auto Regressive and Moving Average are some of the most powerful, yet simple models for time-series forecasting. They can be used individually or together as ARMA. There are many other variations as well. We will use these models to forecast time-series in this worksheet

## Task

Cryptocurrency is all the rage now and it uses the very exciting technology behind blockchain. People even claim it to be revolutionary. But if you have invested in cryptocurrencies, you know how volatile these cryptocurrencies really are! People have become billionaires by investing in crypto, and others have lost all their money on crypto. The most recent instance of this volatility was seen during the Terra Luna crash. Find more info about that [here](#) and [here](#) if you are interested.

Your task is to effectively forecast the prices of **DogeCoin**, a crypto that started as a meme but now is a crypto that people actually invest in. DogeCoin prices however, are affected even by a single tweet by Elon Musk. The image below tweeted by Elon Musk shot up the prices of DogeCoin by 200%!

You have been provided with the daily prices of DogeCoin from 15-08-2021 to 15-08-2022 a period of 1 year (365 days) in the file `doge.csv`

Please download the data from this [Github repo](#)

## Data Dictionary

Date - Date on which price was recorded

Price - Price of DogeCoin on a particular day

## Data Ingestion and Preprocessing

- Read the file into a Pandas DataFrame object

```
import pandas as pd
df = pd.read_csv('doge.csv')
```

```
df.head()
```

	Date	Price
0	2021-08-15	0.348722
1	2021-08-16	0.349838
2	2021-08-17	0.345208
3	2021-08-18	0.331844
4	2021-08-19	0.321622

## Prerequisites

- Set up a new conda env or use an existing one that has `jupyter-notebook` and `ipykernel` installed (Conda envs come with these by default) [Reference](#)
- Instead, you can also use a python venv and install `ipykernel` manually (We instead suggest using conda instead for easy setup) [Reference](#)
- Install the `statsmodels` package either in your Conda environment or Python venv. Refer to [the installation guide](#)

```
!pip install ipykernel
```

```
!pip install statsmodels
```

## Points

The problems in this worksheet are for a total of 10 points with each problem having a different weightage.

- Problem 0: 0.5 points
- Problem 1: 1.5 point
- Problem 2: 2 points
- Problem 3: 1 points
- Problem 4: 2 point
- Problem 5: 1 point
- Problem 6: 1 points

## HINTS FOR ALL PROBLEMS:

- Consider using `inplace=True` or assign it to new DataFrame, when using pandas transformations. If none of these are done, the DataFrame will remain the same
- Search for functions in the `statsmodels` [documentation](#)

## Problem 0 (0.5 point)

- Set the index of DataFrame to the Date column to make it a time series

```
df.set_index('Date', inplace=True)
df
```

	Price
Date	
2021-08-15	0.348722

```

2021-08-16    0.349838
2021-08-17    0.345208
2021-08-18    0.331844
2021-08-19    0.321622
...
2022-08-11    0.072978
2022-08-12    0.073563
2022-08-13    0.073670
2022-08-14    0.079436
2022-08-15    0.082882

```


```
[366 rows x 1 columns]
```

### Problem 1 (1.5 point)

- Plot the time-series. Analyze the stationarity from the time-series. Provide reasoning for stationarity/non-stationarity based on visual inspection of time-series plot (0.5 point)
- Plot the ACF plot of the Time series (upto 50 lags). Analyze the stationarity from ACF plot and provide reasoning (Hint: look at functions in statsmodels package) (1 Point)

```
df['Price'].plot()
```


```
<matplotlib.axes._subplots.AxesSubplot at 0x7f28fa353650>
```



```
#Plotting the ACF
```

```
import statsmodels.api as sm
import matplotlib.pyplot as plt
```

```
sm.graphics.tsa.plot_acf(df.values.squeeze(), lags=50)
plt.show()
```



The time series plot for the given dataset doesn't follow the rules of stationarity. That is 1) The means values of Price is at different values of time (date) are not constant. 2) The variances of Price on different time periods(dates) are not constant. 3) The covariances in prices at time  $t$  and  $t-k$  depends on both  $t$  and  $k$ . Also, a gradual decrease in ACF, rather than a sharp decrease/ cut off to 0 is another sign of nonstationarity. Hence, the given time series is non-stationary.

### Problem 2 (2 point)

- Run Augmented Dickey Fuller Test. Analyze whether the time-series is stationary, based on ADF results (1 point)

Hint: Use the `print_adf_results` function below to print the results of the ADF function cleanly after obtaining results from the library function. Pass the results from library function to `print_adf_results` function

```
def print_adf_results(adf_result):
    print('ADF Statistic: %f' % adf_result[0])
    print('p-value: %f' % adf_result[1])
    print('Critical Values:')
    for key, value in adf_result[4].items():
        print('\t%s: %.3f' % (key, value))
```

- If not stationary, apply appropriate transformations. Run the ADF test again to show stationarity after transformation (1 Point)

Hint: diff and dropna. Assign the DataFrame after transformation to a new DataFrame with name transformed\_df

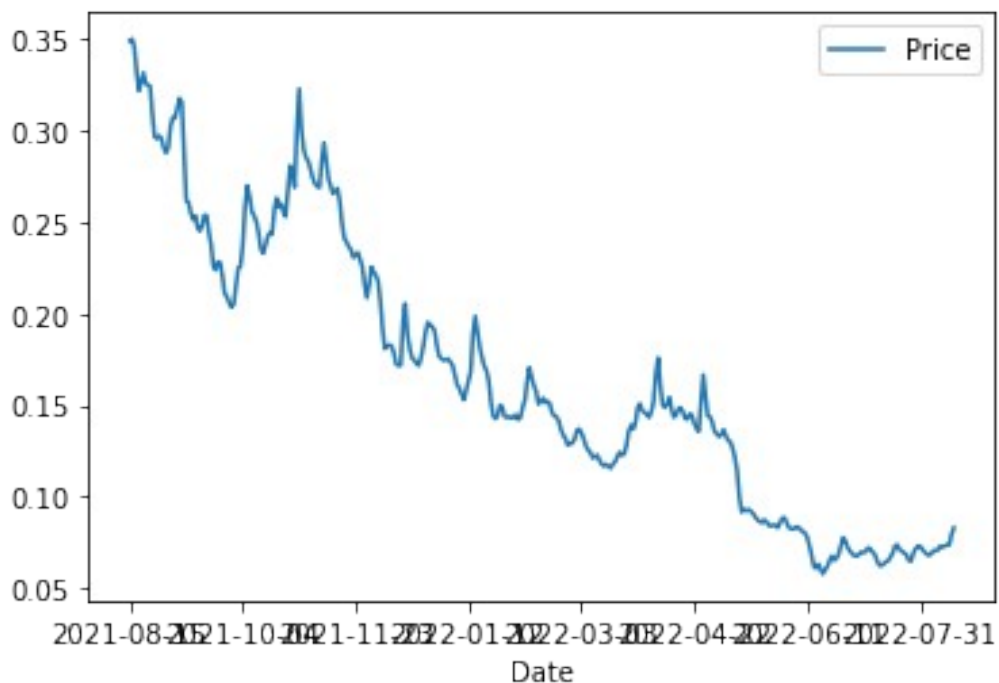
Here the null and alternate hypothesis are as follows:  $H_0$  = Time series is non-stationary.  
 $H_a$  = Time series is stationary.

```
from statsmodels.tsa.stattools import adfuller
adf_result = adfuller(df)
print('ADF Statistic: %f' % adf_result[0])
print('p-value: %f' % adf_result[1])
print('Critical Values:')
for key, value in adf_result[4].items():
    print('\t%s: %.3f' % (key, value))
```

```
ADF Statistic: -1.558935
p-value: 0.504182
Critical Values:
    1%: -3.449
    5%: -2.870
   10%: -2.571
```

```
df.plot()
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f28fa2aec0>
```



The ADF statistics for the original dataframe are as give above. A p-value = 0.5 is obtained which signifies fail to reject  $H_0$ . Even the plot shows a non-stationary/varying wave. Thus, we can conclude that the time series is non-stationary.

```
transformed_df = df.diff().dropna()
trans_adf = adfuller(transformed_df)
print_adf_results(trans_adf)
```

ADF Statistic: -5.593446

p-value: 0.000001

Critical Values:


1%: -3.449

5%: -2.870

10%: -2.571

```
transformed_df.plot()
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f28fa238410>



After applying the transformation of a single order difference and dropping all NA values, we obtain a new dataframe called `transformed_df`. The ADF statistics and the p-value(0.000001) indicate rejection of  $H_0$  and conclude that the new time series is stationary.


### Problem 3 (1 point)

- Plot both ACF and PACF plot. From these select optimal parameters for the ARIMA(p,q) model

Hint: Negative values that are significantly outside the Confidence interval are considered significant too.


Hint:  $p+q = 3$

```
import statsmodels.api as sm
import matplotlib.pyplot as plt
sm.graphics.tsa.plot_acf(transformed_df, lags=50)
plt.show()
```



As seen from the ACF plot, there are 2 values that are significantly outside the confidence interval compared to all of the other values. Also the auto-correlation function decreases exponentially after 2 instances. Hence value of  $p$  can be chosen to be 2.

```
sm.graphics.tsa.plot_pacf(transformed_df, lags=50)
plt.show()
```



As seen from the above PACF plot, the partial auto-correlation is  $> 1.96/\sqrt{n}$  for the order 1. Thus, the value for  $q$  should be 1.

#### Problem 4 (2 point)

- Write a function to forecast values using only AR( $p$ ) model (2 Points)  
Only use pandas functions and Linear Regression from sklearn. [LR documentation](#)

Hint: Create  $p$  new columns in a new DataFrame that is a copy of transformed\_df  
Each new column has lagged value of Price. Price\_t-i (From Price\_t-1 upto Price\_t-p)

Look at the shift function in pandas to create these new columns [Link](#)

```
### Adding columns for lagged values
```

```
arima_df = transformed_df.copy()
```

```
## AR terms
```

```
p = 2
```

```
# Creating p new columns, for p lagged values
```

```
for i in range(1,p+1):
```

```
    arima_df[f'Price_t-{i}'] = arima_df['Price'].shift(i)
```

```
arima_df.dropna(inplace=True)
```

```
arima_df
```

	Price	Price_t-1	Price_t-2
Date			
2021-08-18	-0.013363	-0.004630	0.001116



2021-08-19	-0.010222	-0.013363	-0.004630
2021-08-20	0.004498	-0.010222	-0.013363
2021-08-21	0.005169	0.004498	-0.010222
2021-08-22	-0.005841	0.005169	0.004498
...	...	...	...
2022-08-11	0.000380	-0.000374	0.001993
2022-08-12	0.000585	0.000380	-0.000374
2022-08-13	0.000107	0.000585	0.000380
2022-08-14	0.005767	0.000107	0.000585
2022-08-15	0.003446	0.005767	0.000107

[363 rows x 3 columns]

Hint:

- **Seperate into X\_train and y\_train for linear regression**
- We know that AR(p) is linear regression with p lagged values, and we have created p new columns with the p lagged values
- X\_train is training input that consists of the columns Price\_t-1 upto Price\_t-p (p columns in total)
- y\_train is the training output (truth values) of the Price, i.e the Price column (Only 1 column)

```
X_train = arima_df[['Price_t-1', 'Price_t-2']].values
y_train = arima_df['Price'].values
```

- Set up the Linear Regression between X\_train and y\_train [LR documentation](#)

Name the LinearRegression() object lr

```
from sklearn.linear_model import LinearRegression
lr = LinearRegression()
lr.fit(X_train, y_train)
```

```
LinearRegression()
```

```
lr.coef_
```

```
array([ 0.57292128, -0.39148166])
```

*# Adding new column with predictions using the LR coefficients. The LR Coefficients are Alpha values or AR coefficients*

```
arima_df['AR_Prediction'] = X_train.dot(lr.coef_.T) + lr.intercept_
```

```
arima_df
```

Date	Price	Price_t-1	Price_t-2	AR_Prediction
2021-08-18	-0.013363	-0.004630	0.001116	-0.003682
2021-08-19	-0.010222	-0.013363	-0.004630	-0.006436

```


2021-08-20  0.004498  -0.010222  -0.013363      -0.001218
2021-08-21  0.005169   0.004498  -0.010222      0.005986
2021-08-22 -0.005841   0.005169   0.004498      0.000608
...
2022-08-11  0.000380  -0.000374   0.001993     -0.001587
2022-08-12  0.000585   0.000380  -0.000374     -0.000228
2022-08-13  0.000107   0.000585   0.000380     -0.000406
2022-08-14  0.005767   0.000107   0.000585     -0.000760
2022-08-15  0.003446   0.005767   0.000107      0.002669

```

[363 rows x 4 columns]

```
arima_df.plot(y=['AR_Prediction', 'Price'])
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f28fa097990>



Once you get predictions like this using AR you would have to, undifference the predictions (which are differenced), but we will not deal with that here. For some hints on how to undifference the data to get actual predictions look [here](#)

### Problem 5 (1 Point)

Phew! Just handling AR(2) manually required us to difference, apply regression, undifference. Let's make all of this much easier with a simple library function

- **Use the ARIMA function using parameters picked to forecast values (1 point)**

Hint: Look at ARIMA function the `statsmodels`. Pass the `p, d, q` inferred from the previous tasks

We **DO NOT** need to pass the transformed\_df to the ARIMA function.  
 Pass the original df as input to ARIMA function, with the d value inferred when  
 Transforming the df to make it stationary  
 The ARIMA function automatically performs the differencing based on the d value passed  
 Store the .fit() results in an object named res

```
import statsmodels.api as sm
model = sm.tsa.arima.ARIMA(df, order=(2,1,1))
res = model.fit()
res.summary()
```

```
/usr/local/lib/python3.7/dist-packages/statsmodels/tsa/base/
tsa_model.py:527: ValueWarning: No frequency information was provided,
so inferred frequency D will be used.
% freq, ValueWarning)
/usr/local/lib/python3.7/dist-packages/statsmodels/tsa/base/tsa_model.
py:527: ValueWarning: No frequency information was provided, so
inferred frequency D will be used.
% freq, ValueWarning)
/usr/local/lib/python3.7/dist-packages/statsmodels/tsa/base/tsa_model.
py:527: ValueWarning: No frequency information was provided, so
inferred frequency D will be used.
% freq, ValueWarning)
/usr/local/lib/python3.7/dist-packages/statsmodels/base/model.py:568:
ConvergenceWarning: Maximum Likelihood optimization failed to
converge. Check mle_retvals
ConvergenceWarning)
```

```
<class 'statsmodels.iolib.summary.Summary'>
"""
```

## SARIMAX Results

```
=====
=====
Dep. Variable:          Price    No. Observations:
366
Model:                ARIMA(2, 1, 1)    Log Likelihood
1440.153
Date:                Fri, 07 Oct 2022    AIC
2872.305
Time:                14:14:00    BIC
2856.706
Sample:                08-15-2021    HQIC
2866.106
- 08-15-2022

Covariance Type:                opg

=====
=====
```

	coef	std err	z	P> z	[0.025
0.975]					
-----					
-----					
ar.L1	-0.0540	0.048	-1.137	0.256	-0.147
0.039					
ar.L2	-0.1093	0.052	-2.115	0.034	-0.211
-0.008					
ma.L1	0.9396	0.021	43.763	0.000	0.898
0.982					
sigma2	2.201e-05	7.48e-07	29.428	0.000	2.05e-05
2.35e-05					

```

=====
=====
Ljung-Box (L1) (Q):          0.01  Jarque-Bera (JB):
3047.97
Prob(Q):                    0.91  Prob(JB):
0.00
Heteroskedasticity (H):     0.15  Skew:
1.73
Prob(H) (two-sided):        0.00  Kurtosis:
16.73
=====
=====

```

Warnings:

```

[1] Covariance matrix calculated using the outer product of gradients
(complex-step).
"""

```


```

## TODO: Use ARIMA function

# # Making predictions and plotting
df['Predictions'] = res.predict(0, len(df)-1)
df.plot()


<matplotlib.axes._subplots.AxesSubplot at 0x7f28fad6c8d0>

```



```
# Forecast for 20 future dates after training data ends
res.forecast(20).plot()
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f28f9f73d50>
```



### Problem 6 (1 point)

- Evaluate the ARIMA model using Ljung Box test. Based on p-value infer if the Model shows lack of fit

Hint: Pass the `res.resid` (Residuals of the ARIMA model) as input the Ljung-Box Text. Pass `lags=[10]`. Set `return_df=True` For inference, refer back to the Null and Alternate Hypotheses of Ljung-Box test. (If p value high, Null Hypothesis is significant)

```
#import statsmodels.api as sm
```

```
sm.stats.acorr_ljungbox(res.resid, lags=[10], return_df=True)
```

	lb_stat	lb_pvalue
10	0.929177	0.999877